

Constant step-size gradient descent vs exact line search

$$||\theta - \eta_*||^2$$

